



Introduction to Deep Learning for Engineers

Spring 25, Introduction to Deep Learning
Jan 14, 2025

Amir Barati Farimani

*Russel V. Trader Associate Professor of Mechanical Engineering and Bio-
Engineering*

Carnegie Mellon University

DALL. E2

“a teddy bear on a skateboard in times square”



Text To Image Generation

Imagen

— A group of teddy bears in suit in a corporate office celebrating the birthday of their friend. There is a pizza cake on the desk.



Imagen

“Hierarchical Text-Conditional Image Generation with CLIP Latents” Ramesh et al. 2022

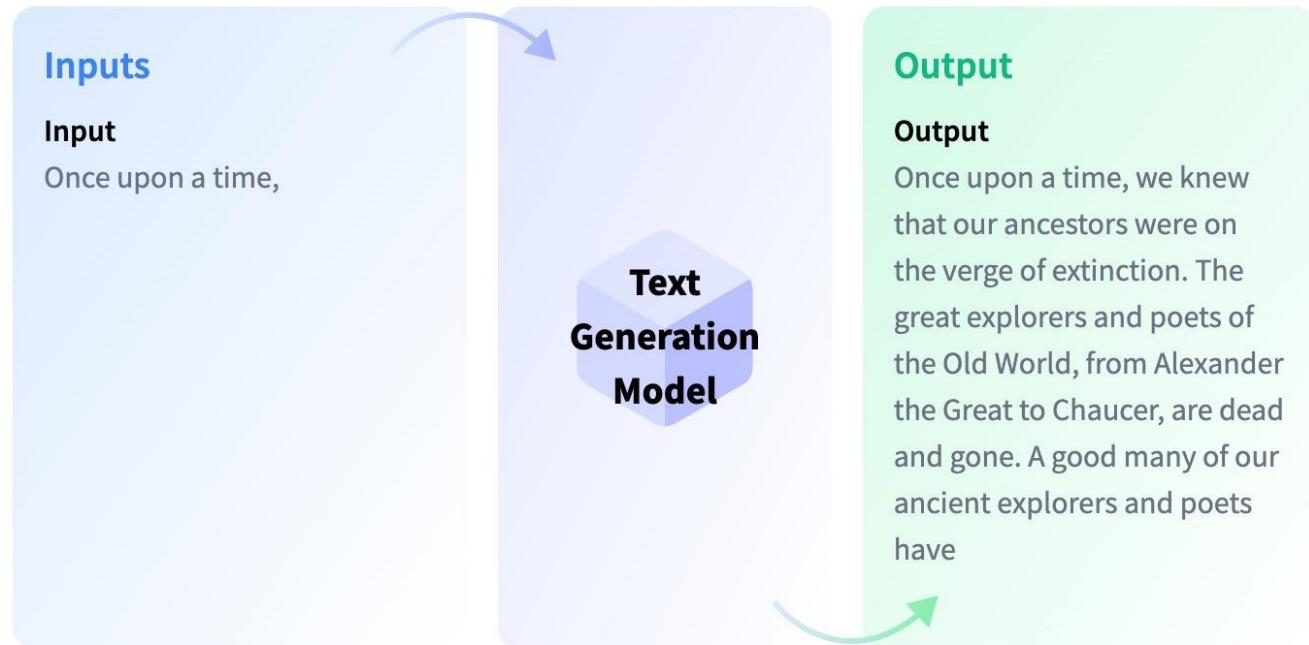
“Photorealistic Tex-to-Image Diffusion Models with Deep Language Understanding” Saharia et al. 2022

Text Generation

Natural Language Generation (NLG)

- Language modeling
- Conditional language modeling

Next word prediction



What do we want computers to do with our data?



Label: “Motorcycle”
Suggest tags
Image search
...



Speech recognition
Music classification
Speaker identification
...



Web search
Anti-spam
Machine translation
...

Computer vision is hard!



Why is this hard?

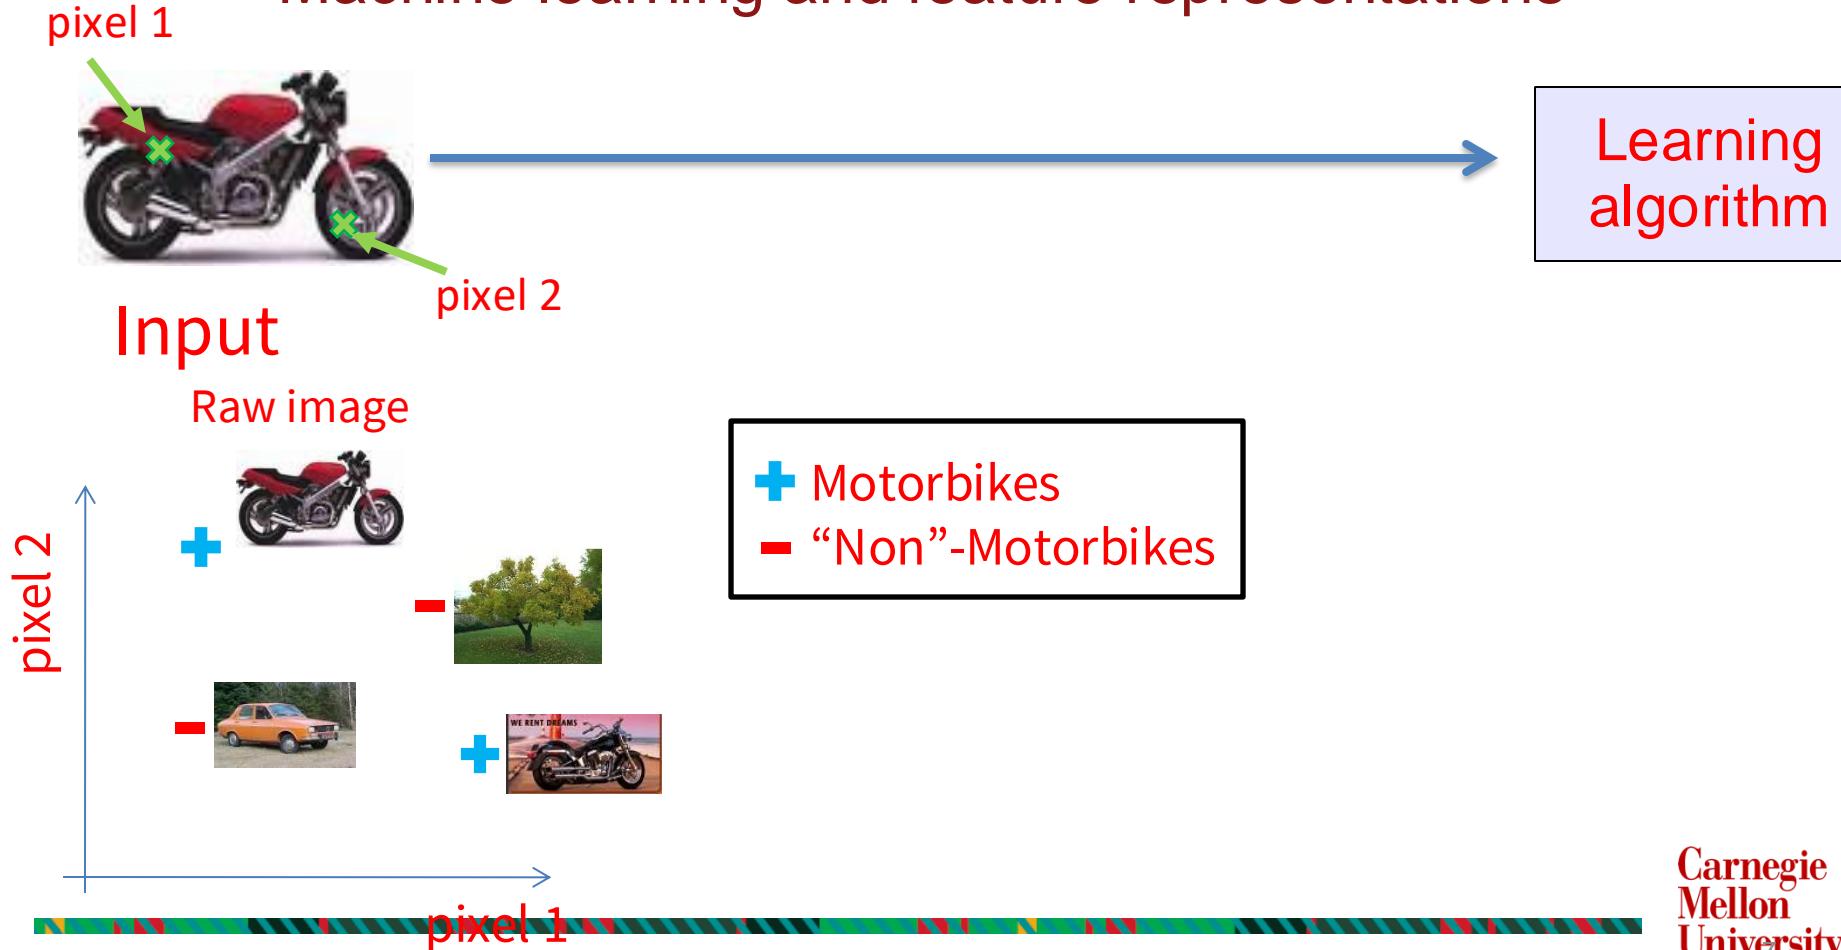
You see this:



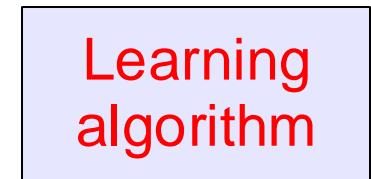
But the camera sees this:

194	210	201	212	199	213	215	195	178	158	182	209
180	189	190	221	209	205	191	167	147	115	129	163
114	126	140	188	176	165	152	140	170	106	78	88
87	103	115	154	143	142	149	153	173	101	57	57
102	112	106	131	122	138	152	147	128	84	58	66
94	95	79	104	105	124	129	113	107	87	69	67
68	71	69	98	89	92	98	95	89	88	76	67
41	56	68	99	63	45	60	82	58	76	75	65
20	43	69	75	56	41	51	73	55	70	63	44
50	50	57	69	75	75	73	74	53	68	59	37
72	59	53	66	84	92	84	74	57	72	63	42
67	61	58	65	75	78	76	73	59	75	69	50

Machine learning and feature representations



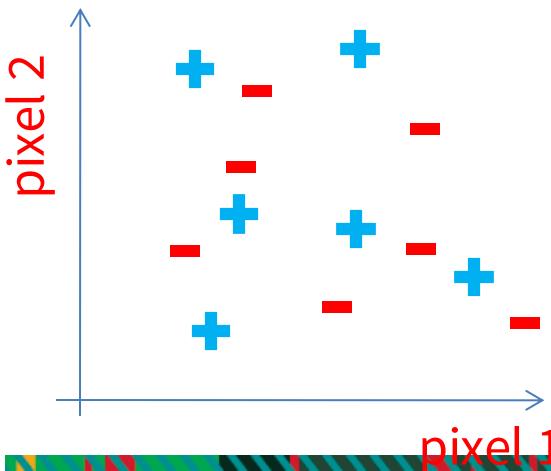
What we want



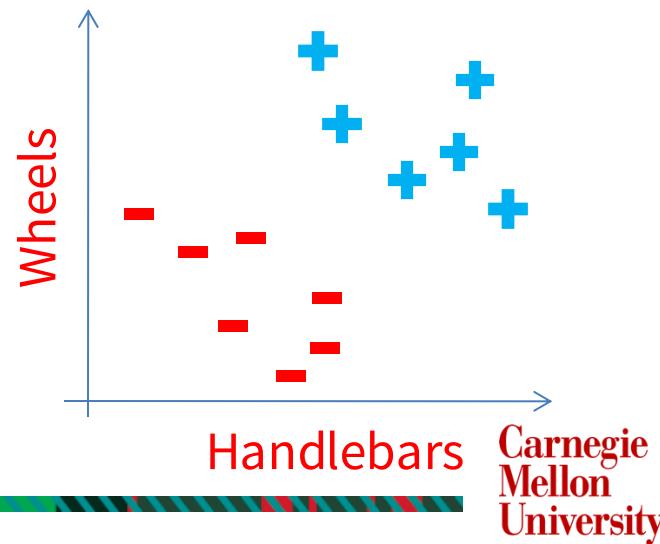
E.g., Does it have Handlebars? Wheels?

Input

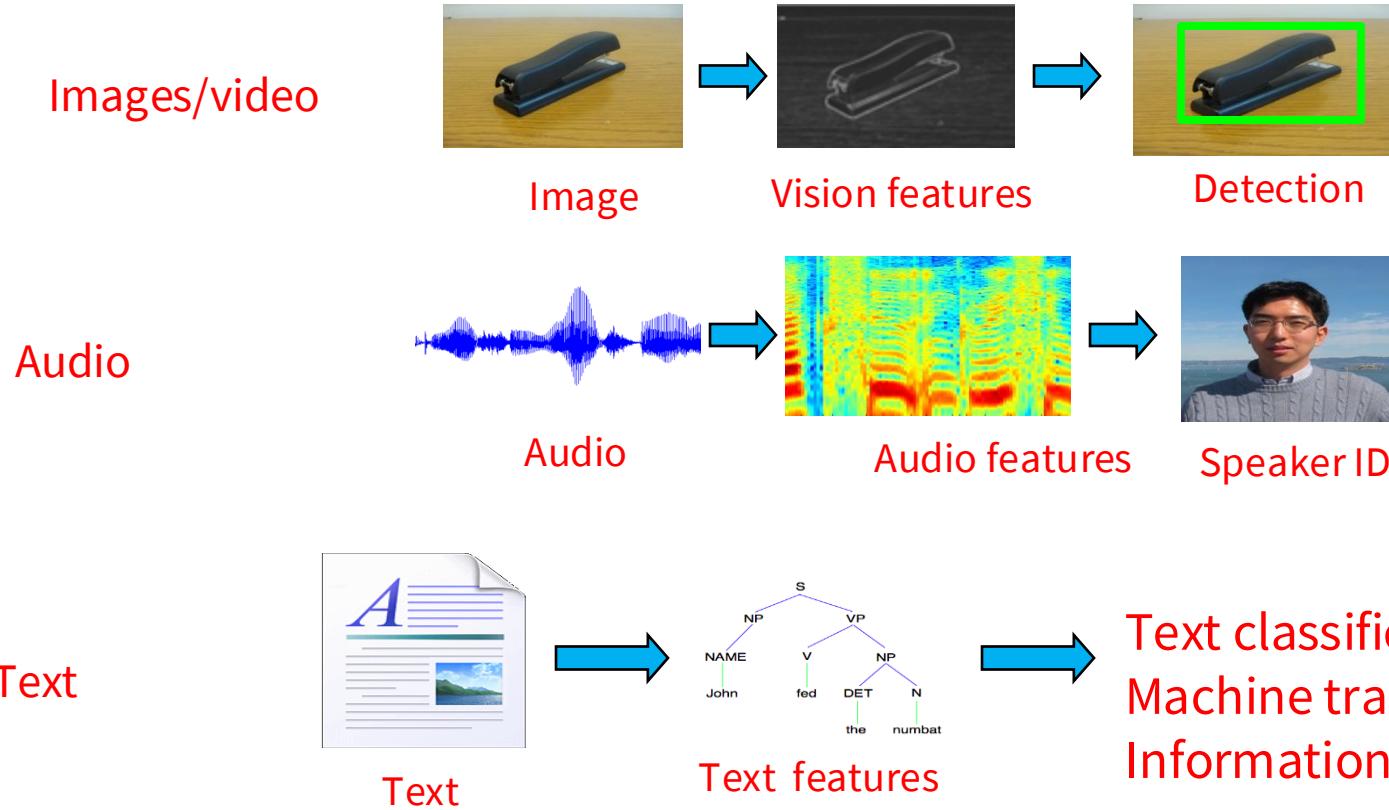
Raw image



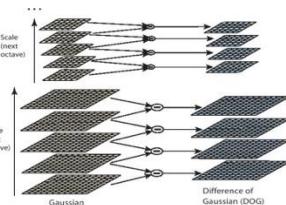
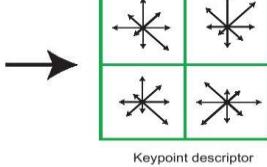
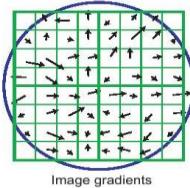
Features



How is computer perception done?

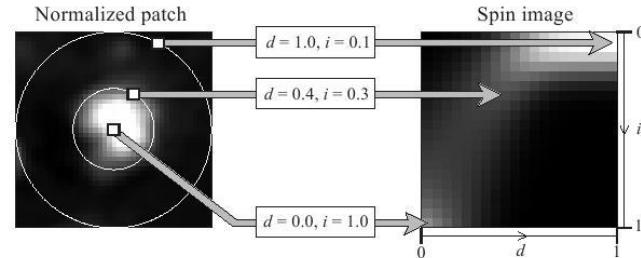
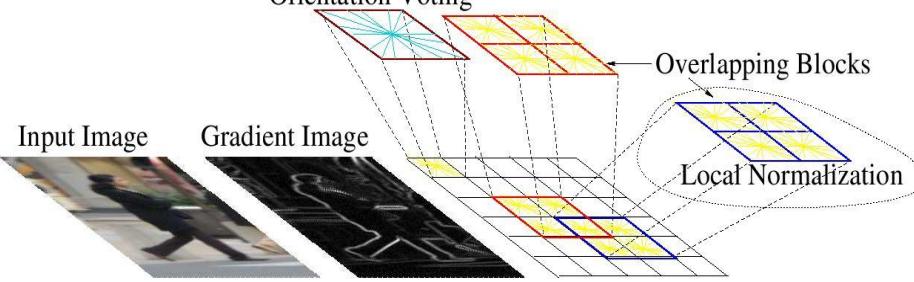


Computer vision features

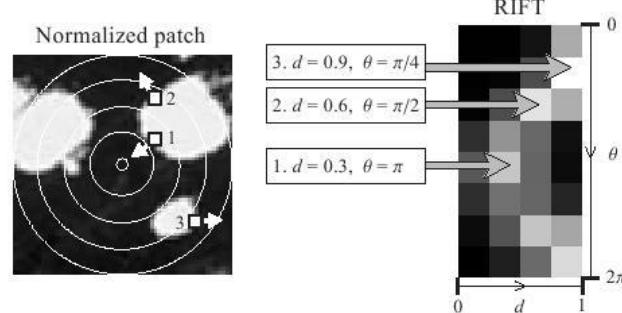


SIFT

Orientation Voting



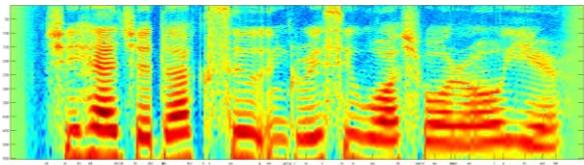
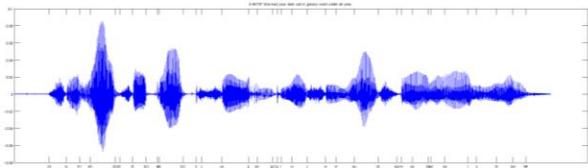
Spin image



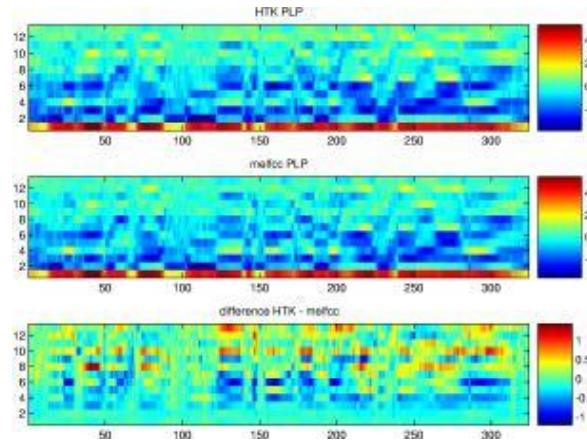
RIFT

HoG

Audio features

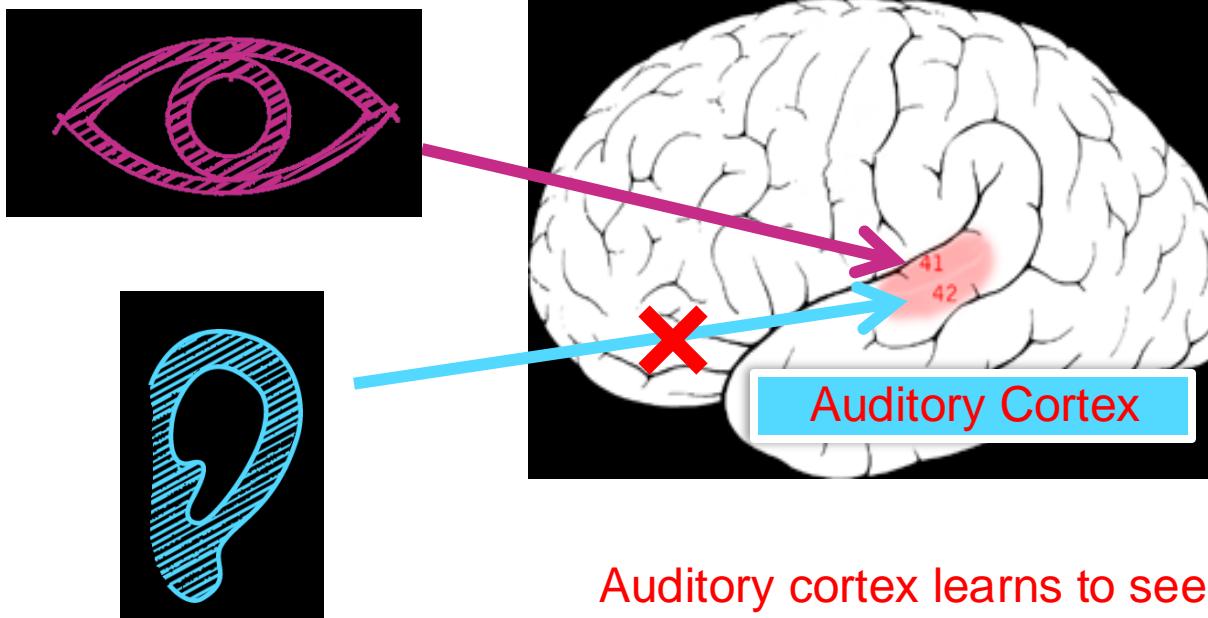


Spectrogram



MFCC

The “one learning algorithm” hypothesis



[Roe et al., 1992]

We Need a Machine Learning Universal Representation Learning Method

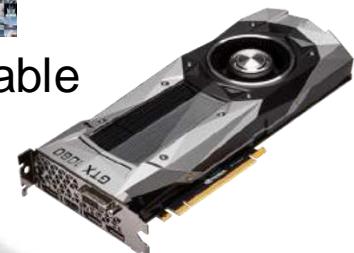
The Answer: Deep Learning

Why Now?

- **Big Data:** Appearance of big datasets and challenges such as ImageNET



- **GPU and Hardware Accelerator:** Massively scalable and Parallelizable

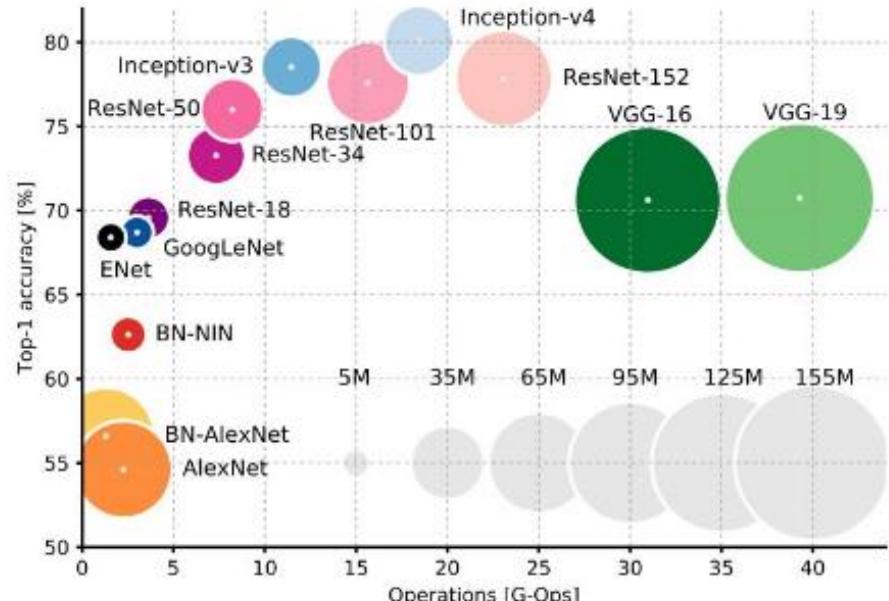
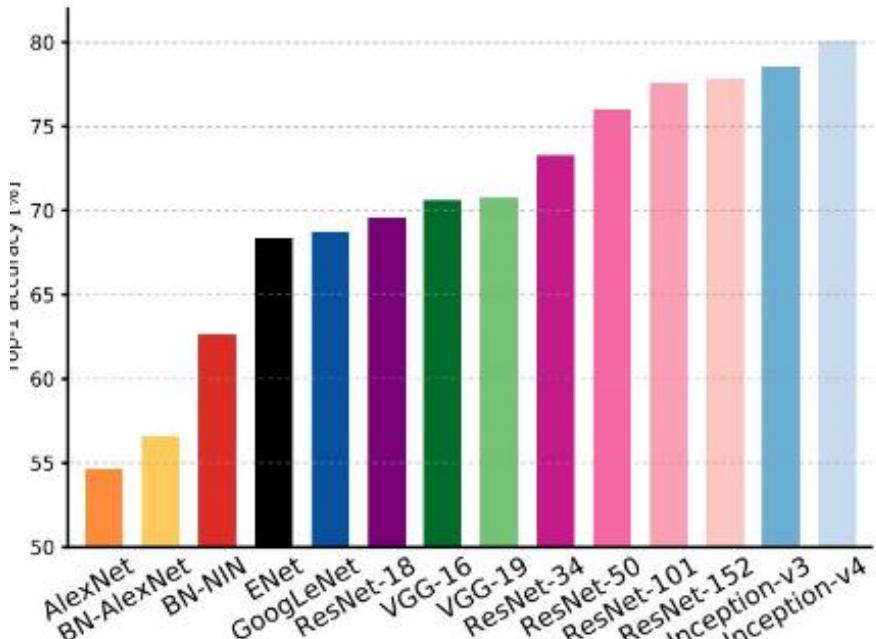


- **Algorithms:** Optimization, Drop out, Batch Normalization,...

- **Software:** Optimization, Drop out, Batch Normalization,...



Why Now?



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

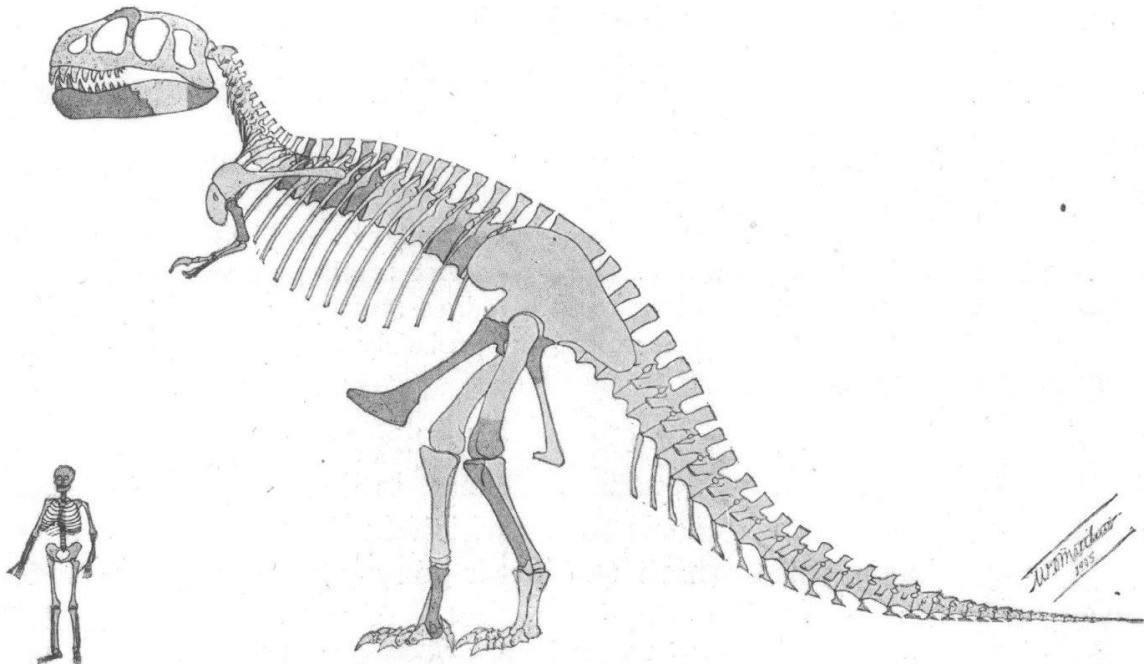
Foundational Models



Carnegie
Mellon
University



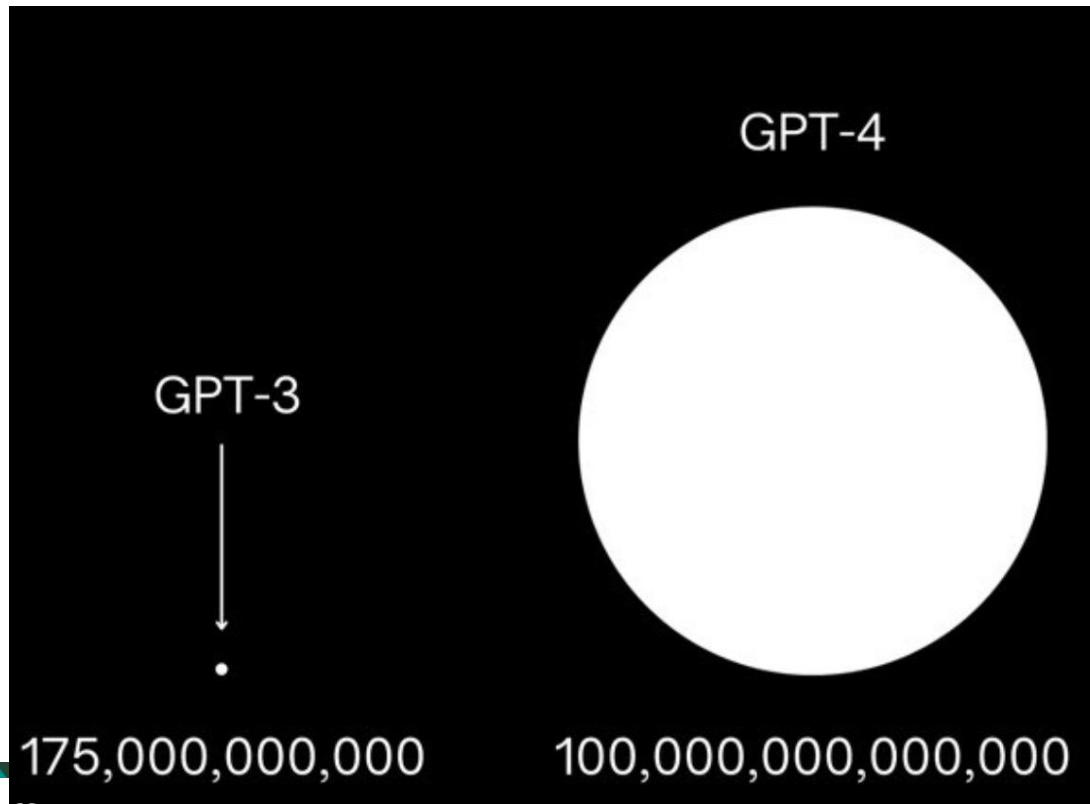
Foundational Models



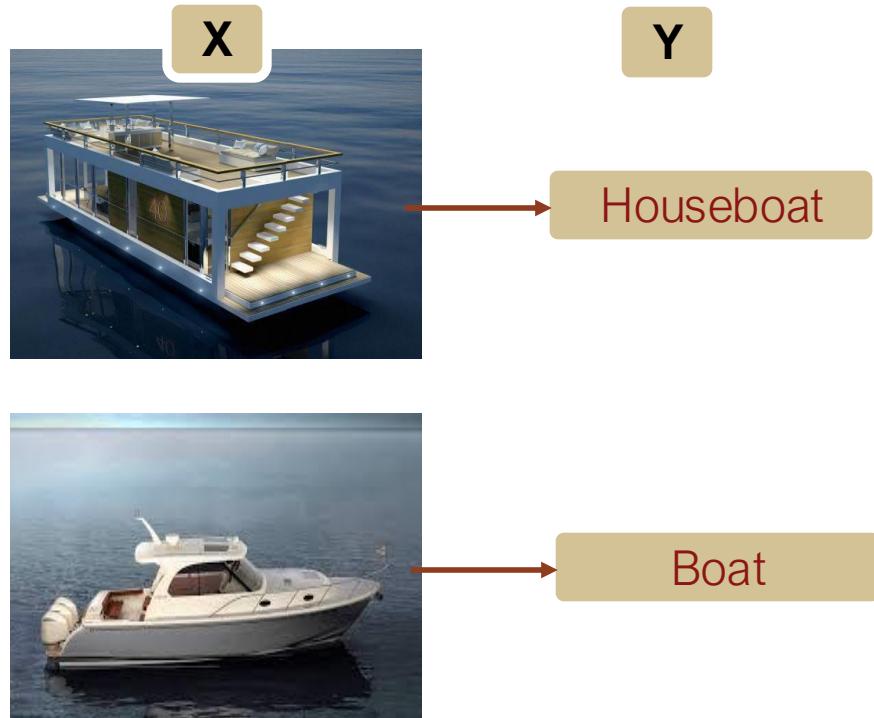
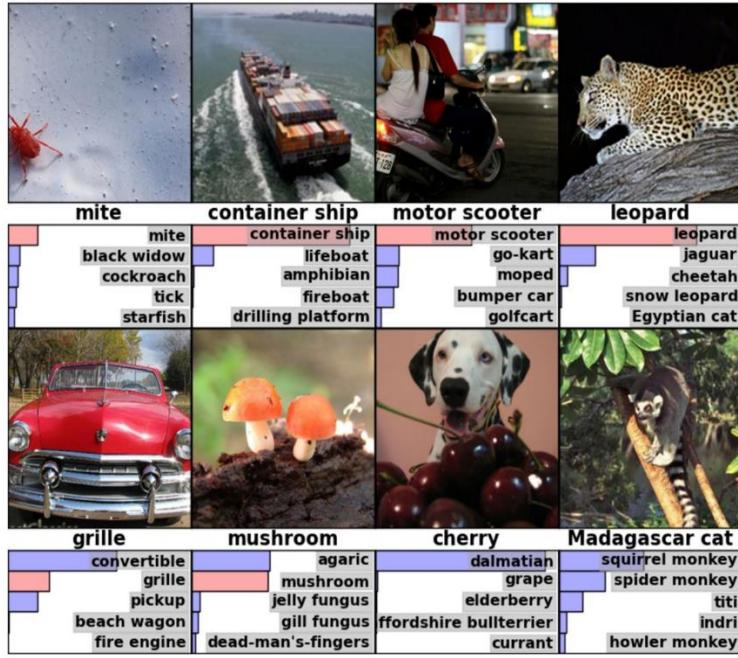
GPT-2
1.5B Parameters

GPT-3
175B Parameters

Foundational Models

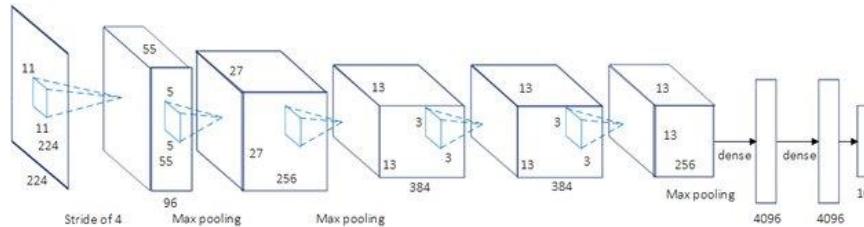
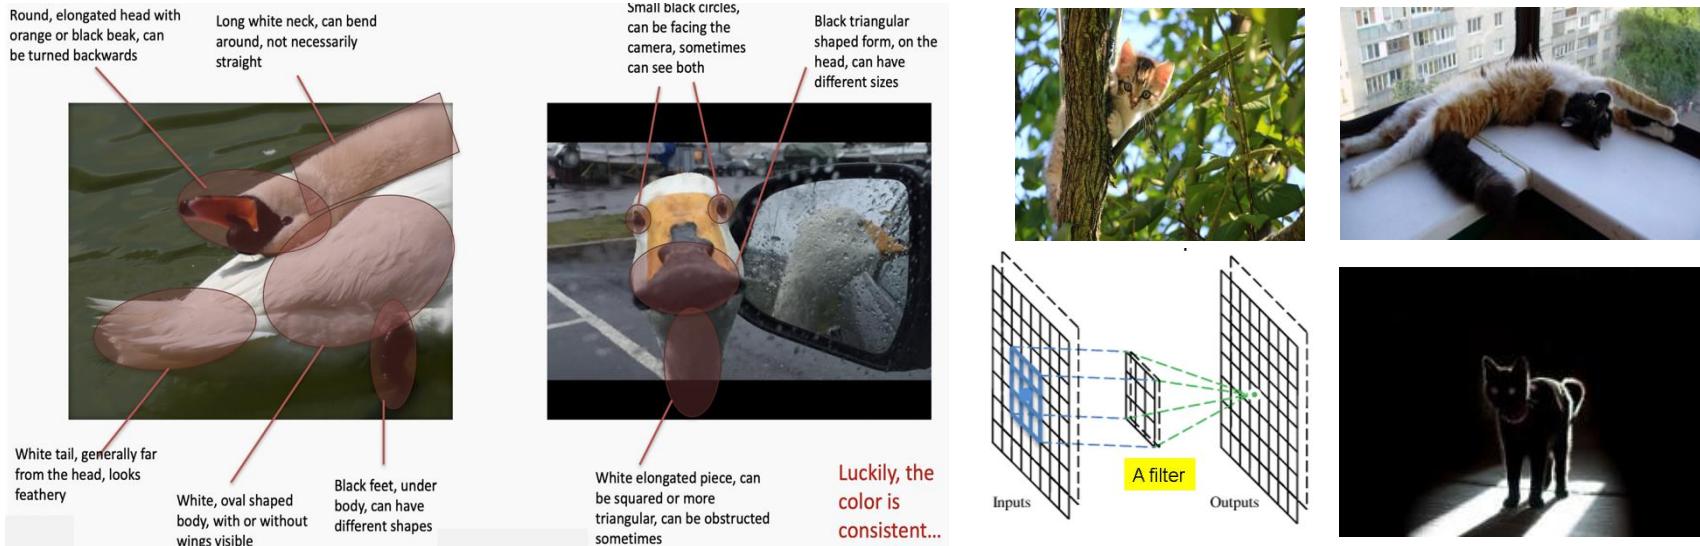


Exuberance of AI



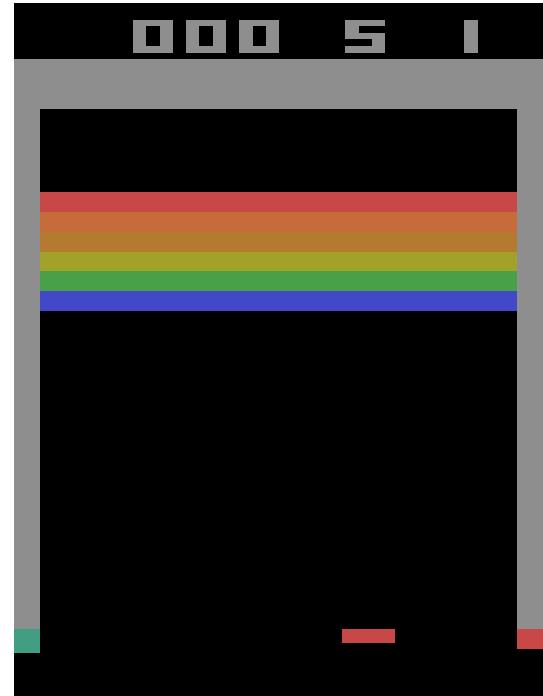
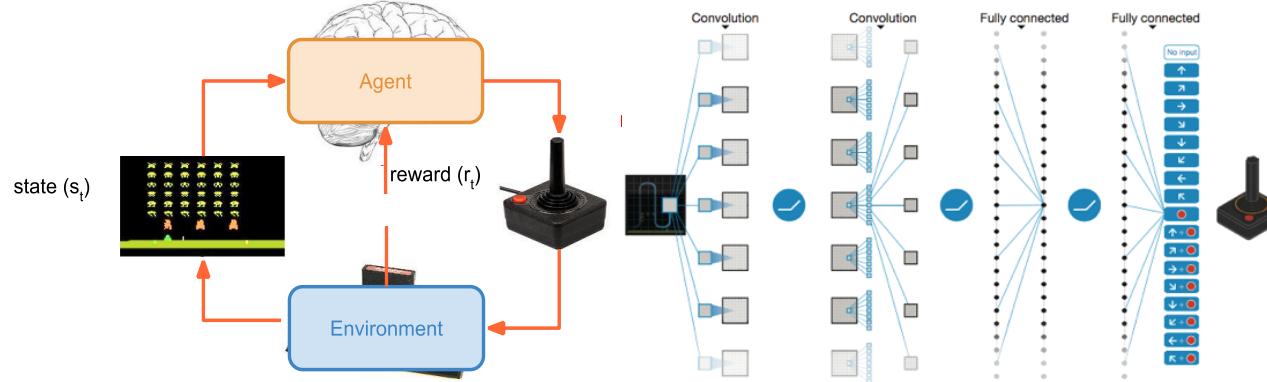
<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>

Revolution 1: Automatic Feature Detection (2012)



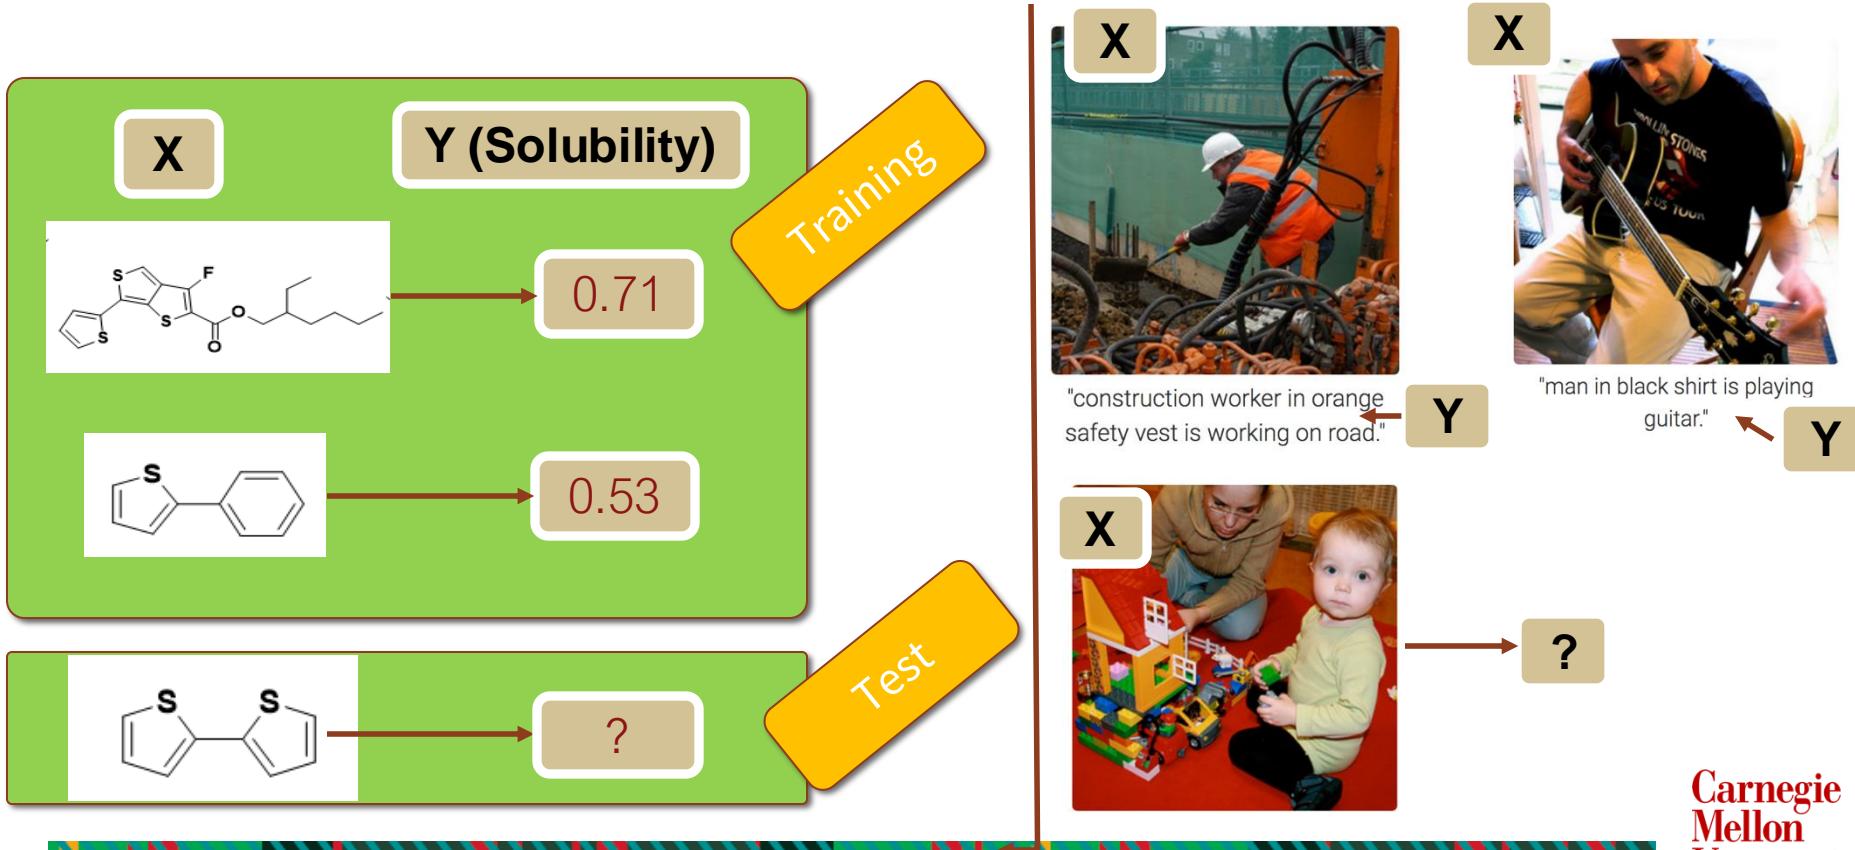
Revolution 2: Deep Reinforcement Learning (2015)

Q-learning + Function Approximation + Deep Network



Human-level control through deep reinforcement learning, Nature, 518, pages, 529–533, (2015)

Why AI and ML?





Godzillium vs. Trumpium:
Some Suggestions to Add
to the Periodic Table



To Protect Against Zika
Virus, Pregnant Women
Are Warned About Latin
American Trips



THE NEW OLD A
F.T.C.'s Lum
Doesn't End
Training De

SCIENCE

Scientists See Promise in Deep-Learning Progr

By JOHN MARKOFF NOV. 23, 2012



Sign in

News

Sport

Weather

Shop

NEWS

Home | Video | World | UK | Business | Tech | Science | Magazi

Forbes / Tech

DEC 29, 2014 @ 11:37 AM 89,471 VIEWS

Tech 2015: Deep Learning And Machine Intelligence Will Eat The World

'Deep learning' technology inspired by human brain

culture business lifestyle fashion environment tech travel

ndroids do dream of electric sheep

...in feedback loop in its image recognition neural network - which

nature

International weekly journal of science

[Home](#) | [News & Comment](#) | [Research](#) | [Careers & Jobs](#) | [Current Issue](#)

[Archive](#) > [Volume 518](#) > [Issue 7540](#) > [News](#) > [Article](#)

NATURE | NEWS

عرب

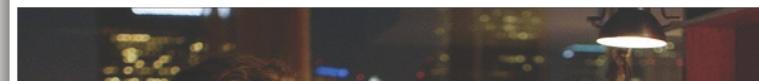
Game-playing software holds lessons for neuroscience

DeepMind computer provides new way to investigate how the brain

Top 20 Stocks for 2016

Google a step closer to developing machines with human-like intelligence

Algorithms developed by Google designed to encode thoughts, could computers with 'common sense' within a decade, says leading AI

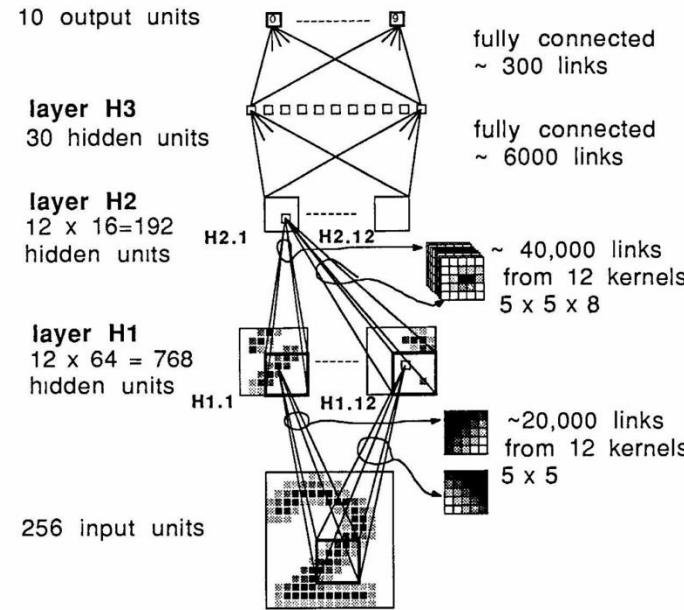


Carnegie
Mellon
University

Milestone 1: Digit Recognition

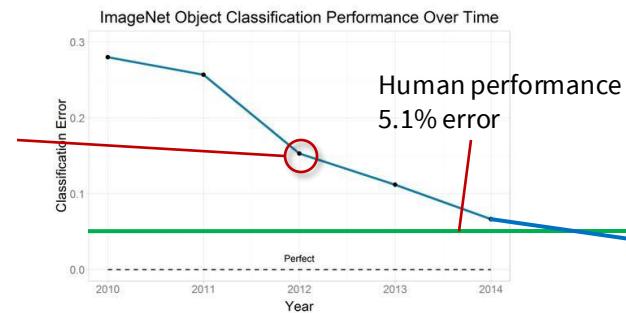
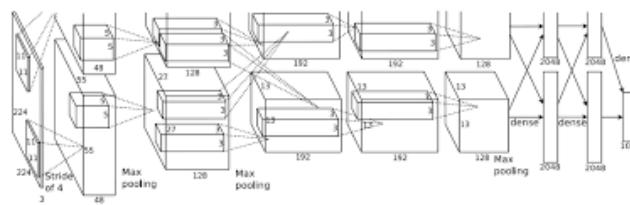
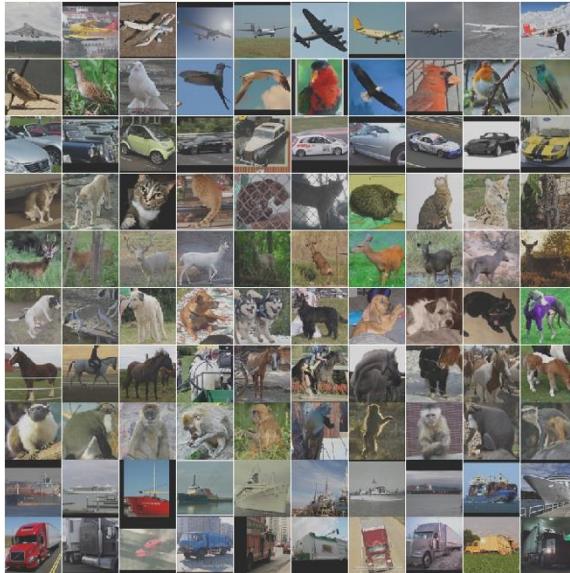
LeNet 1989: recognize zip codes, Yann LeCun, Bernhard Boser and others, ran live in US postal service

80322-4129 80206
40004 14310
37872 05753
35502 75216
35460 44209



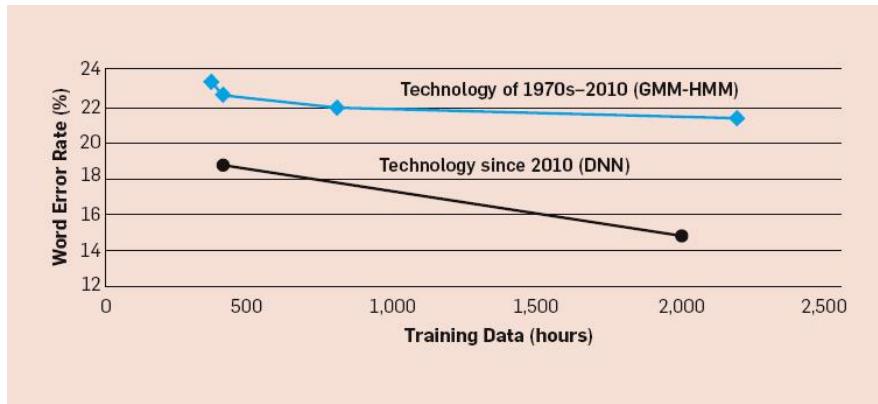
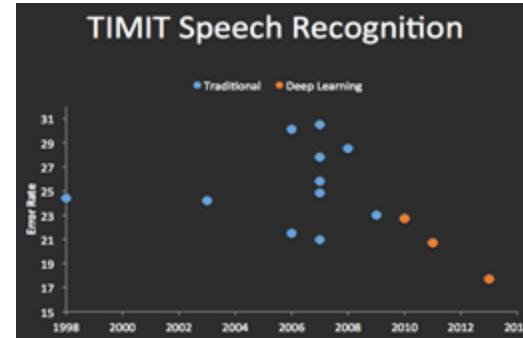
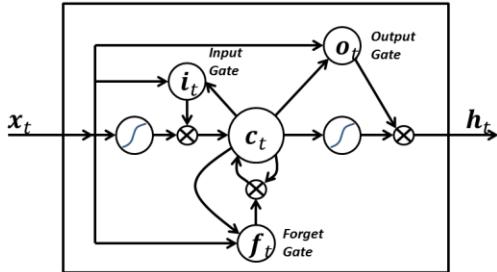
Milestone 2: Image Classification

Convolutional NNs: AlexNet (2012): trained on 200 GB of ImageNet Data



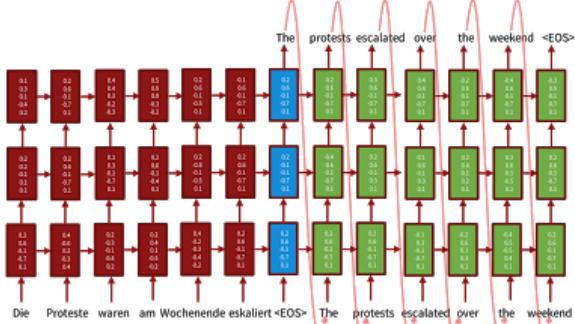
Milestone 3: Speech Recognition

Recurrent Nets: LSTMs (1997):



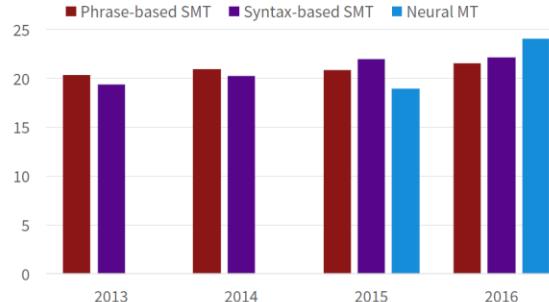
Milestone 4: Machine Translation

Sequence-to-sequence models with LSTMs and attention:



Progress in Machine Translation

[Edinburgh En-De WMT newstest2013 Cased BLEU; NMT 2015 from U. Montréal]

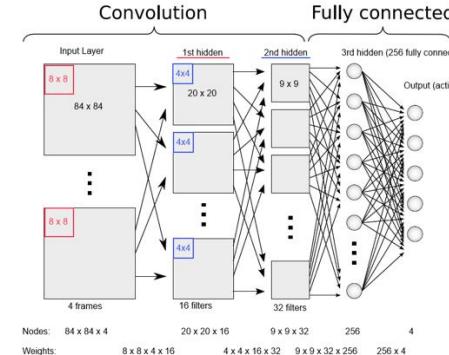
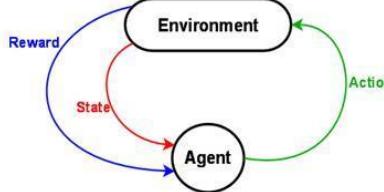


From [Sennrich 2016, http://www.meta-net.eu/events/meta-forum-2016/slides/09_sennrich.pdf]

Source Luong, Cho, Manning ACL Tutorial 2016.

Milestone 5: Deep Reinforcement Learning

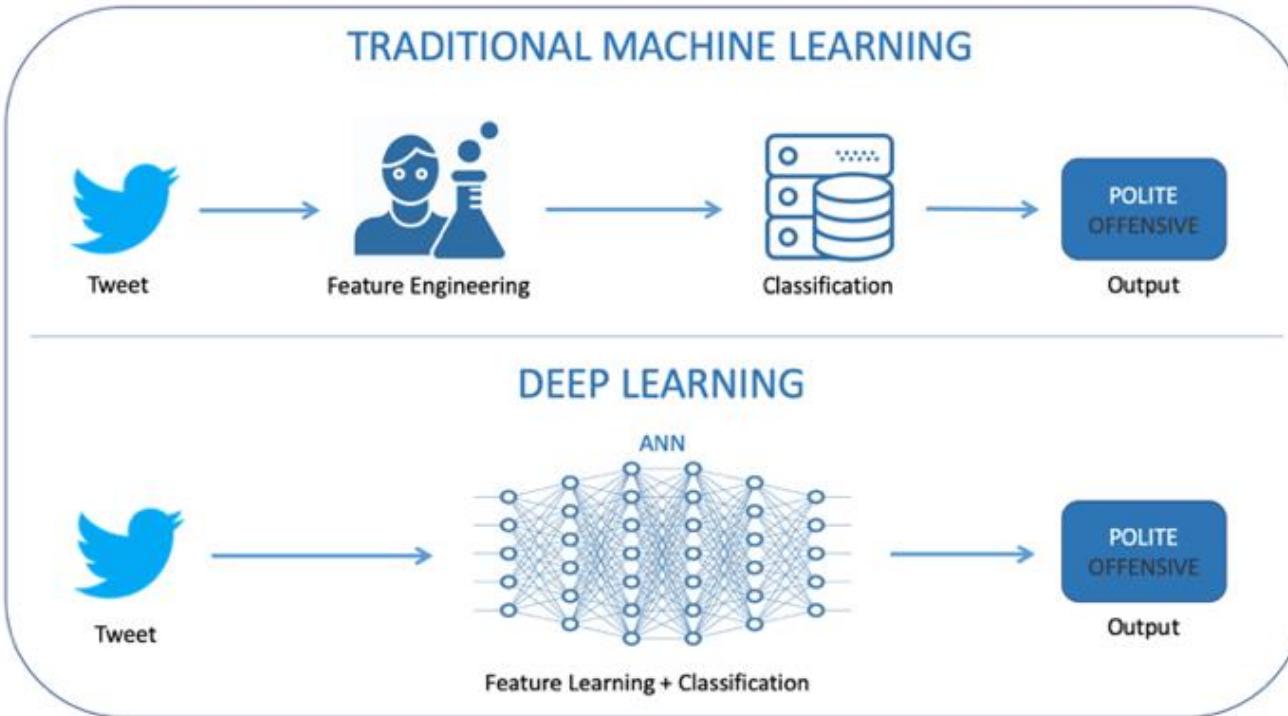
IN 2013, DEEP MIND'S ARCADE PLAYER BESTS HUMAN EXPERT ON SIX ATARI GAMES. ACQUIRED BY GOOGLE IN 2014,.



IN 2016, DEEP MIND'S
ALPHAGO DEFEATS FORMER
WORLD CHAMPION LEE SEDOL



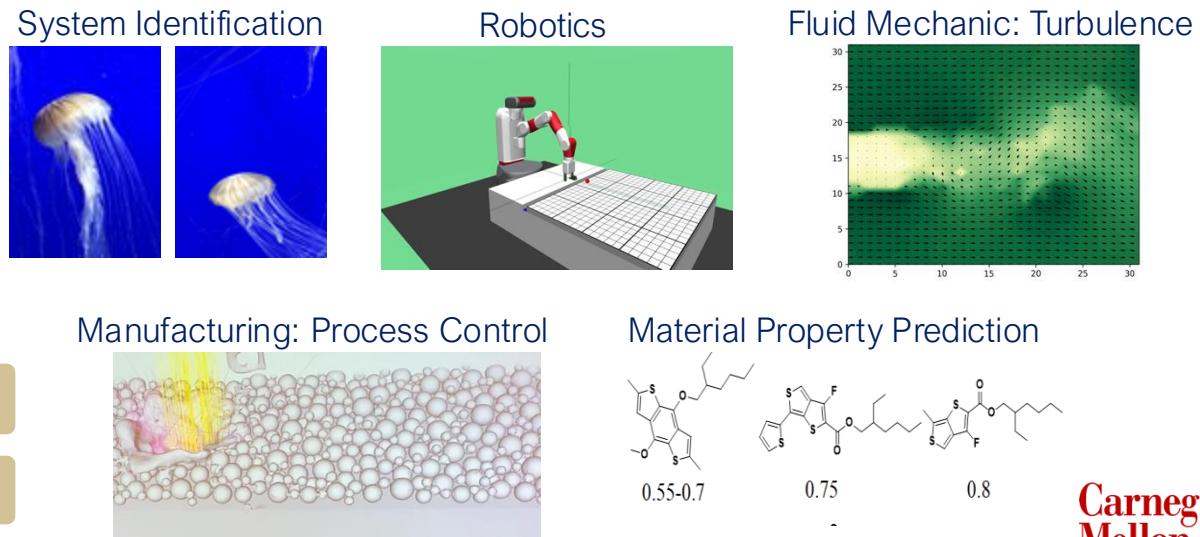
Traditional ML vs Deep Learning



How Engineers Can use AI?

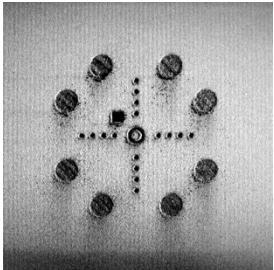
- Complex Phenomena in Engineering that we don't have Mathematical Description
- Where we don't know the Physics, but we want to predict to engineer
- Where we want to speed up computation by surrogate, fast models
- Where we have abundance of data from experiments and want to make a more precise model

Fluid Mechanics
Solid Mechanics
Bio-Engineering
Dynamic and Control
Material Science

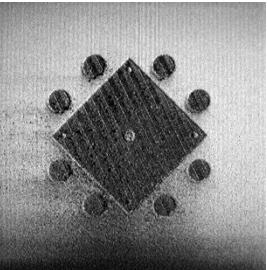


Immediate Use of AI in Engineering

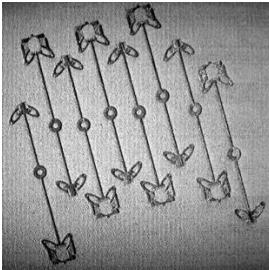
Where images/vision exists



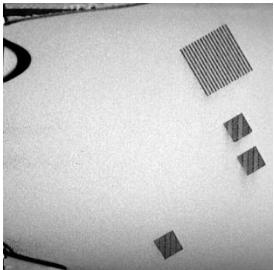
Actual - 2
Predicted - 1



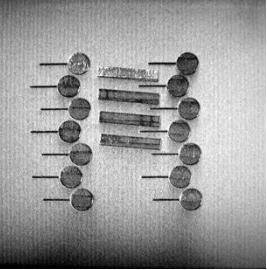
Actual - 2
Predicted - 1



Actual - 0
Predicted - 1



Actual - 1
Predicted - 1



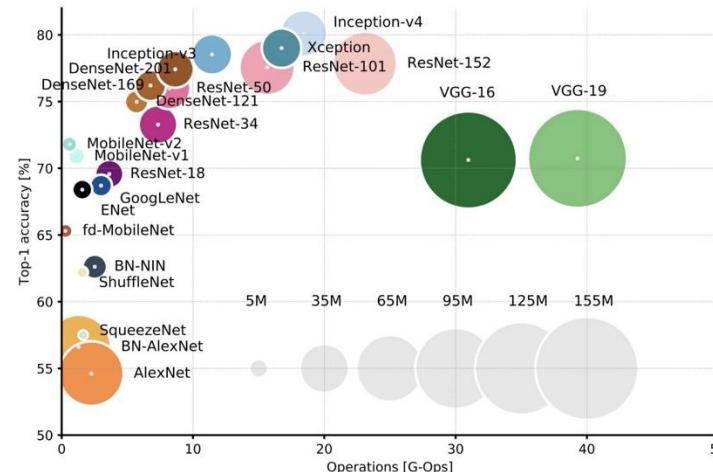
Actual - 2
Predicted - 2



Actual - 1
Predicted - 1

CNN Architecture	% Accuracy
VGG16	80.6
ResNet18	80.9
ResNext50	81.3
Dense Net	81.7
Custom CNN	82.1

- '0' – No spatter
- '1' – Some spatter
- '2' – Too much of spatter



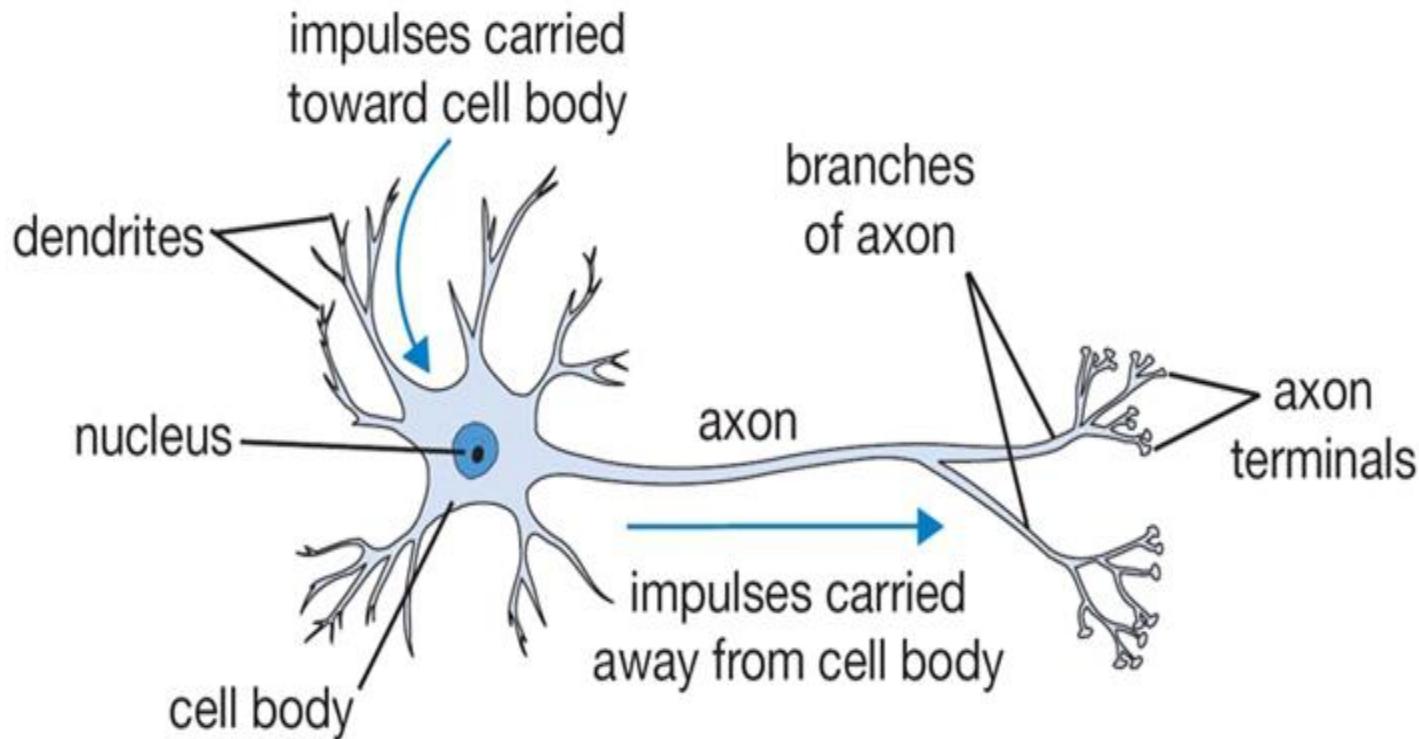
Course Objectives

This course provides an introduction to deep learning. We will learn about the basics of deep neural networks, and their applications to different tasks in engineering. Students will be able to **apply deep learning** to a variety of artificial intelligence tasks pertinent to different engineering problems. Applications of deep learning in Mechanical, chemical, biological, Electrical and material engineering will be discussed.

Biological Neural Network



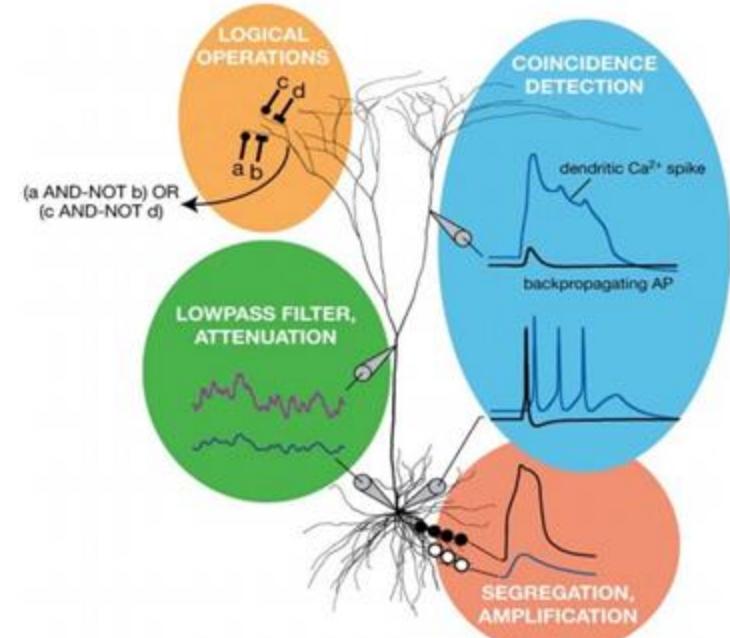
Natural Neuron



Operation of a Biological Neural Network

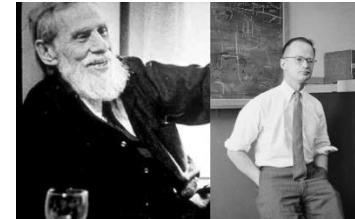
Biological Neurons:

- Many different types
- Dendrites can perform complex non-linear computations
- Synapses are not a single weight but a complex non-linear dynamical system
- Rate code may not be adequate



[Dendritic Computation. London and Häusser]

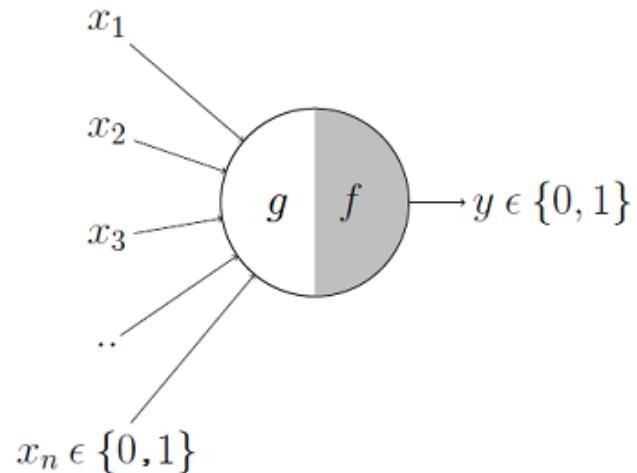
McCulloch-Pitts Neuron



- The first computational model of a neuron was proposed by Warren McCulloch (neuroscientist) and Walter Pitts (logician) in 1943.
- Boolean Functions Using M-P Neuron
- The inputs are all Boolean i.e., $\{0,1\}$
- The outputs are all Boolean

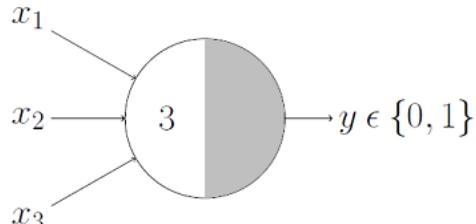
$$g(x_1, x_2, x_3, \dots, x_n) = g(\mathbf{x}) = \sum_{i=1}^n x_i$$

$$\begin{aligned} y = f(g(\mathbf{x})) &= 1 && \text{if } g(\mathbf{x}) \geq \theta \\ &= 0 && \text{if } g(\mathbf{x}) < \theta \end{aligned}$$

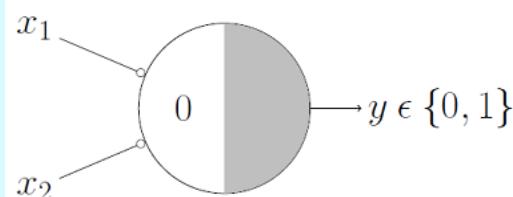


McCulloch-Pitts as Boolean function

AND

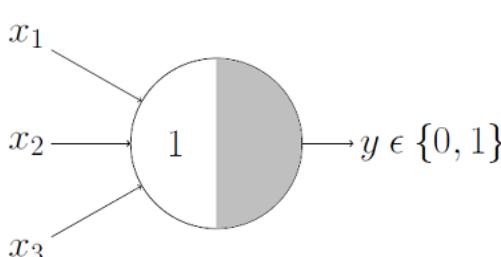


$$g(\mathbf{x}) \geq 3$$



$$g(\mathbf{x}) = 0$$

OR

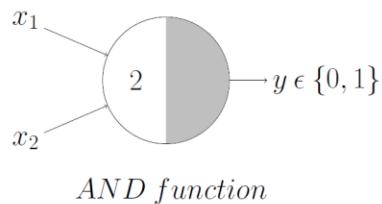


$$g(\mathbf{x}) \geq 1$$

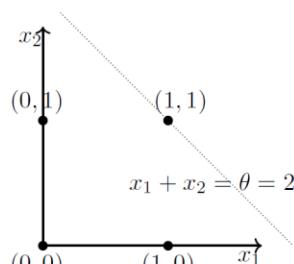
Note: all the inputs and outputs are $\{0, 1\}$

Geometrical Understanding of MP

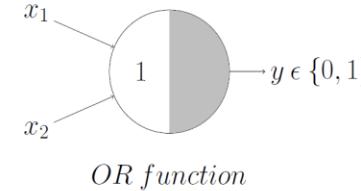
AND with two inputs



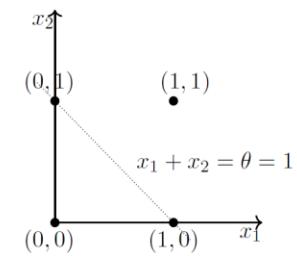
$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 2$$



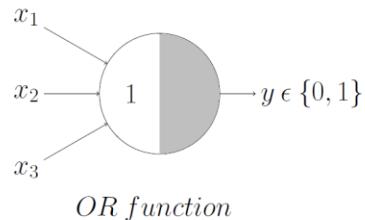
OR with two inputs



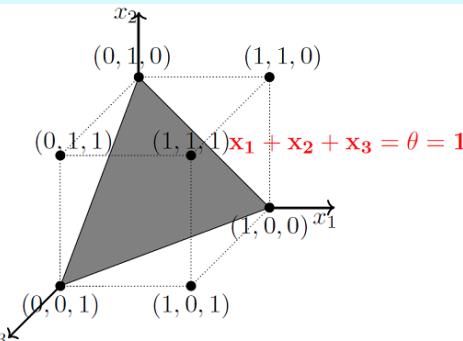
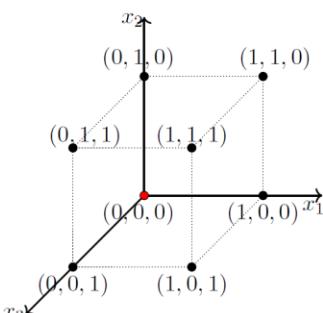
$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 1$$



AND with three inputs



$$x_1 + x_2 + x_3 = \sum_{i=1}^3 x_i \geq 1$$



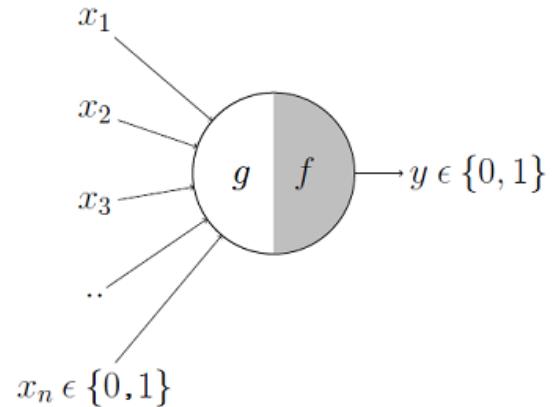
Limitations of M-P Neuron

What about non-boolean (say, real) inputs?

Do we always need to hand code the threshold?

Are all inputs equal? What if we want to assign more importance to some inputs?

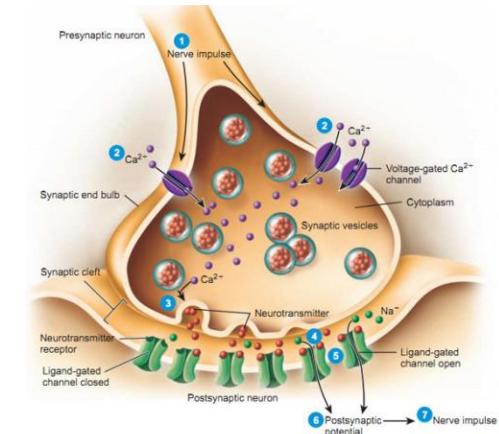
What about functions which are not linearly separable? Say XOR function.



Hebbian Learning

The Hebbian rule was the first learning rule. In 1949 Donald Hebb developed it as learning algorithm of the unsupervised neural network. We can use it to identify how to improve the weights of nodes of a network.

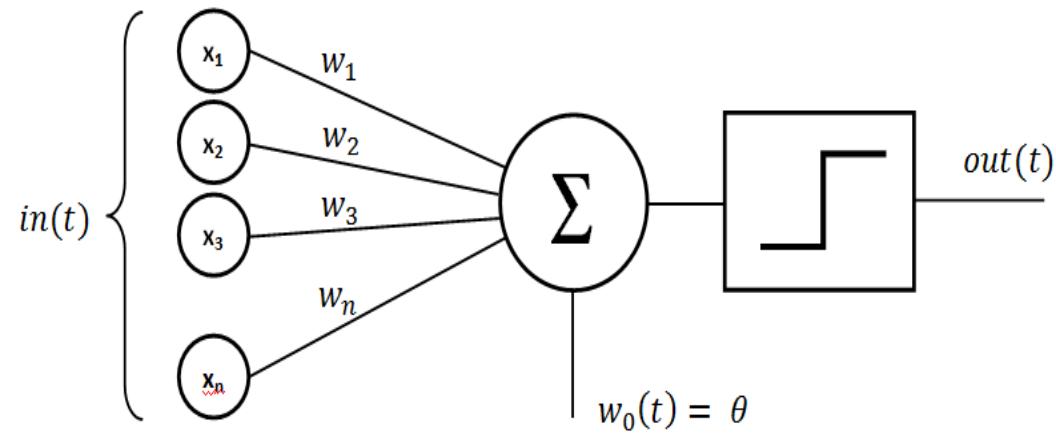
- The Hebb learning rule assumes that:
- If two neighbor neurons activated and deactivated at the same time. Then the weight connecting these neurons should increase.
- For neurons operating in the opposite phase, the weight between them should decrease.
- If there is no signal correlation, the weight should not change.



Perceptron

Frank Rosenblatt

- Psychologist, Logician
- Perceptron Inventor (1958)



Perceptron Algorithm

Algorithm 1: Perceptron Learning Algorithm

Input: Training examples $\{\mathbf{x}_i, y_i\}_{i=1}^m$.

Initialize \mathbf{w} and b randomly.

while not converged **do**

 # # # Loop through the examples.

for $j = 1, m$ **do**

 # # # Compare the true label and the prediction.

$$error = y_j - \sigma(\mathbf{w}^T \mathbf{x}_j + b)$$

 ### If the model wrongly predicts the class, we update the weights and bias.

if $error \neq 0$ **then**

 ### Update the weights.

$$\mathbf{w} = \mathbf{w} + error \times \mathbf{x}_j$$

 ### Update the bias.

$$b = b + error$$

 Test for convergence

Output: Set of weights \mathbf{w} and bias b for the perceptron.

Perceptron

x_1	x_2	OR	
0	0	0	$w_0 + \sum_{i=1}^2 w_i x_i < 0$
1	0	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$
0	1	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$
1	1	1	$w_0 + \sum_{i=1}^2 w_i x_i \geq 0$

$$w_0 + w_1 \cdot 0 + w_2 \cdot 0 < 0 \implies w_0 < 0$$

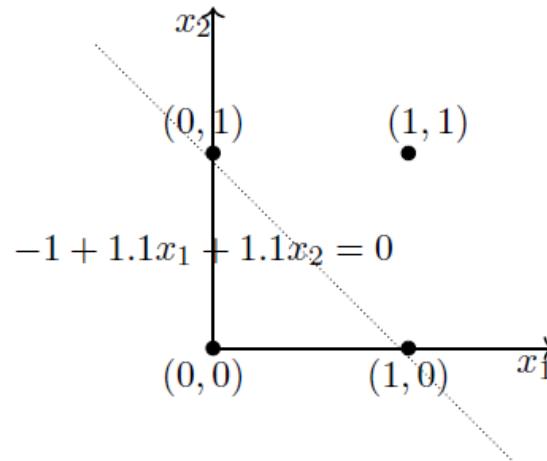
$$w_0 + w_1 \cdot 0 + w_2 \cdot 1 \geq 0 \implies w_2 > -w_0$$

$$w_0 + w_1 \cdot 1 + w_2 \cdot 0 \geq 0 \implies w_1 > -w_0$$

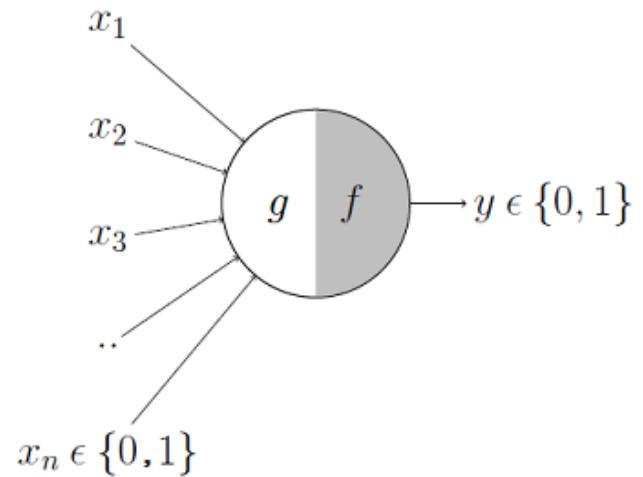
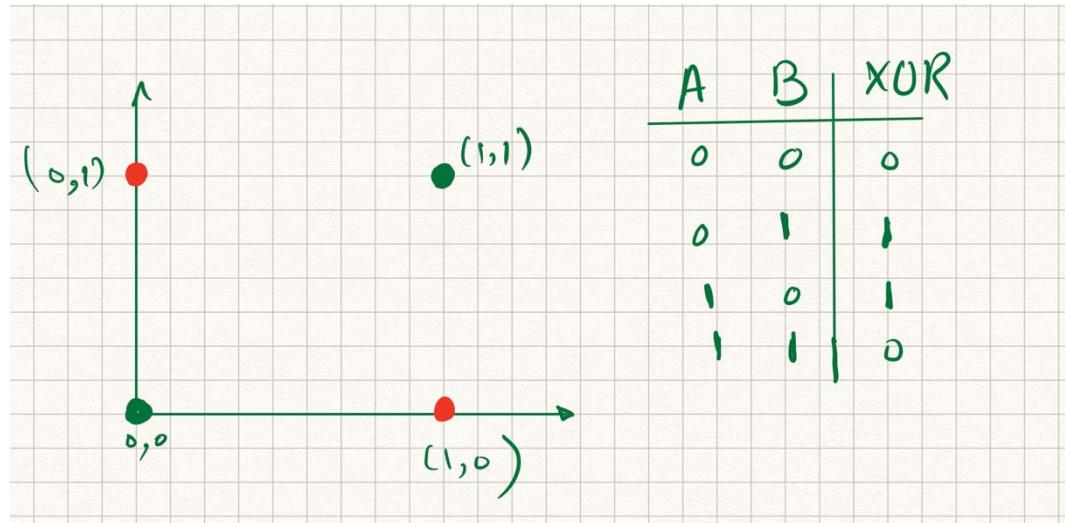
$$w_0 + w_1 \cdot 1 + w_2 \cdot 1 \geq 0 \implies w_1 + w_2 > -w_0$$

One possible solution is

$$w_0 = -1, w_1 = 1.1, w_2 = 1.1$$

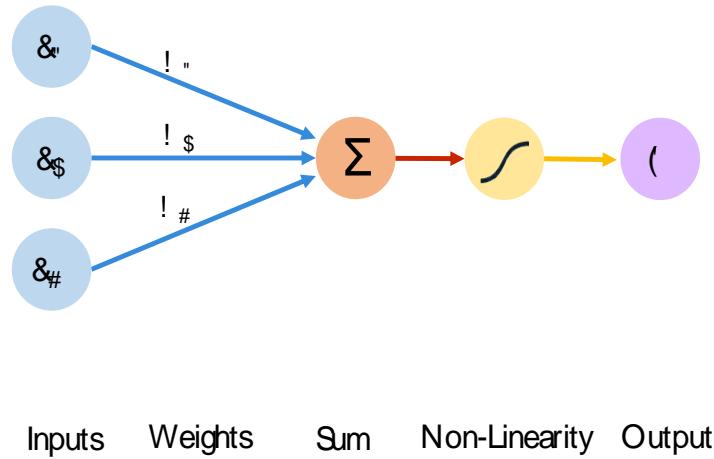
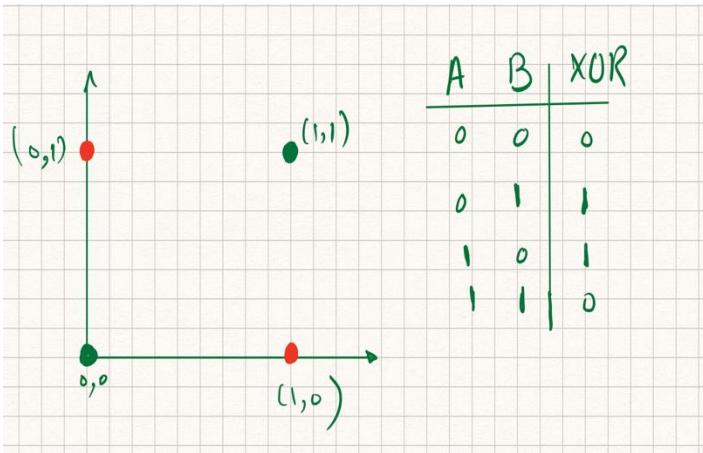


XOR?



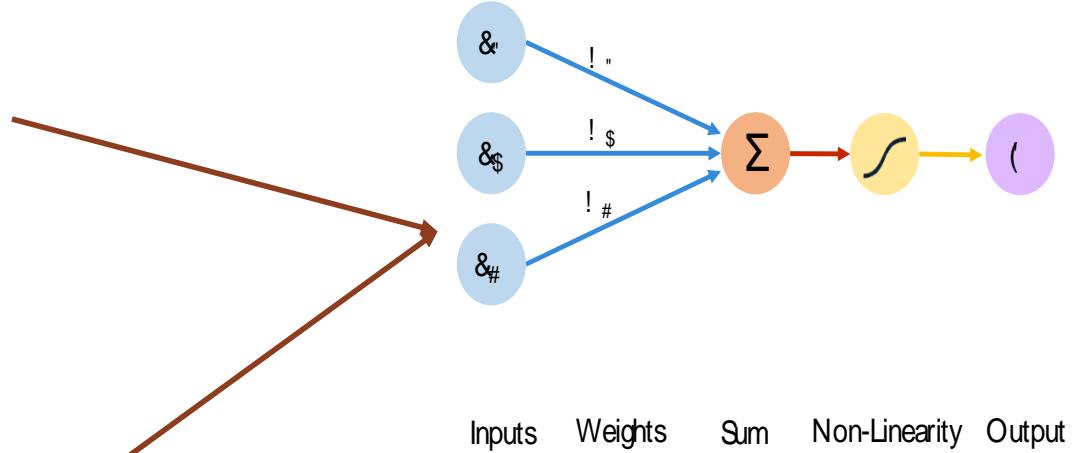
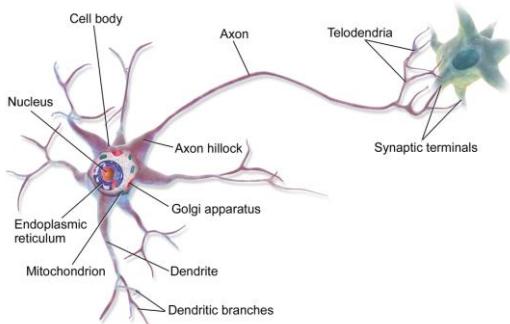
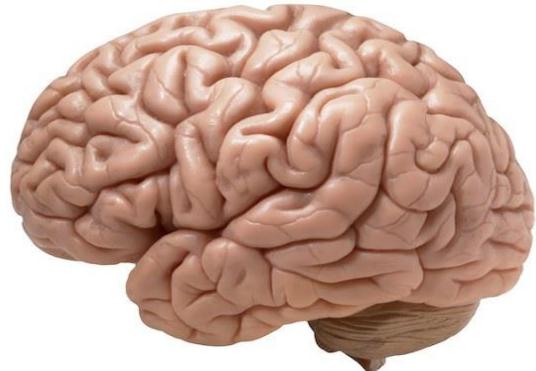
Minsky and Papert, 1968

Perceptron can't do XOR

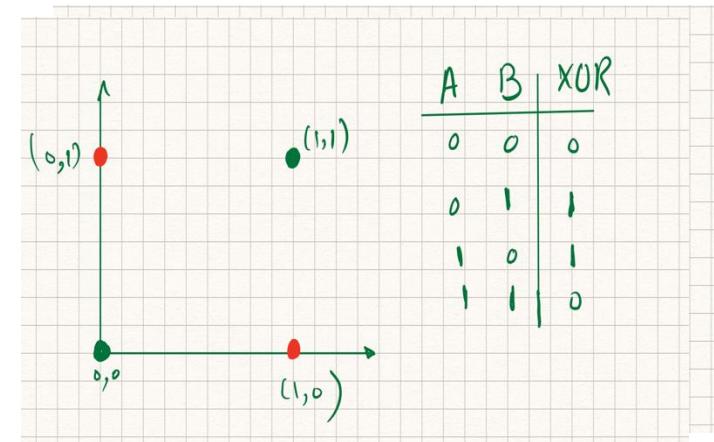
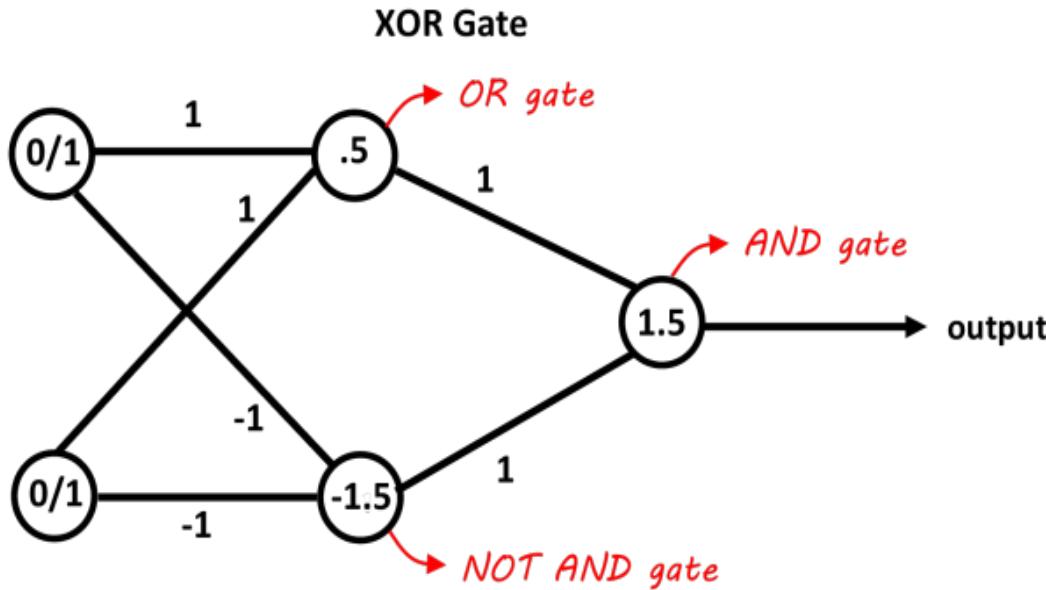


Minsky and Papert in 1969

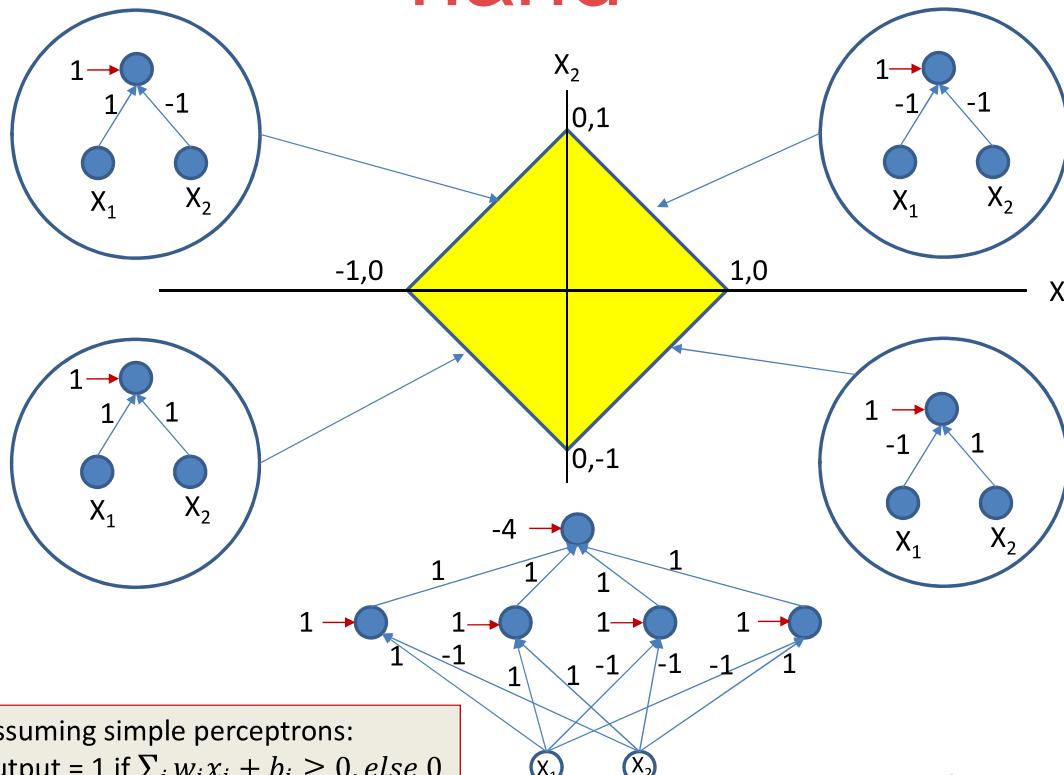
Single Neuron is not ENOUGH



Multilayer Perceptron (MLP) can do XOR



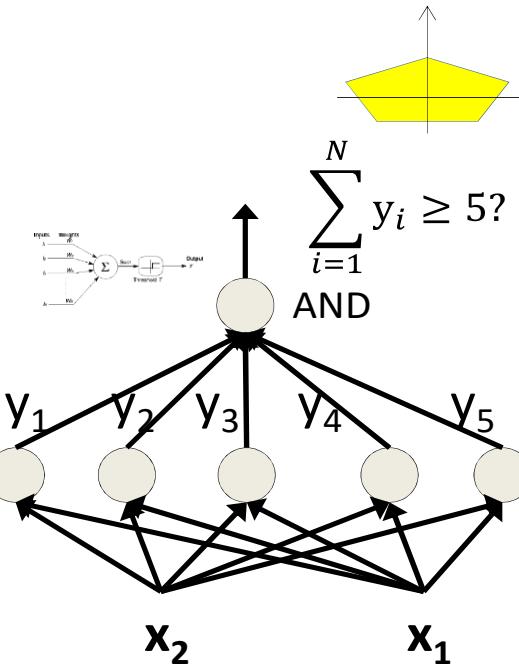
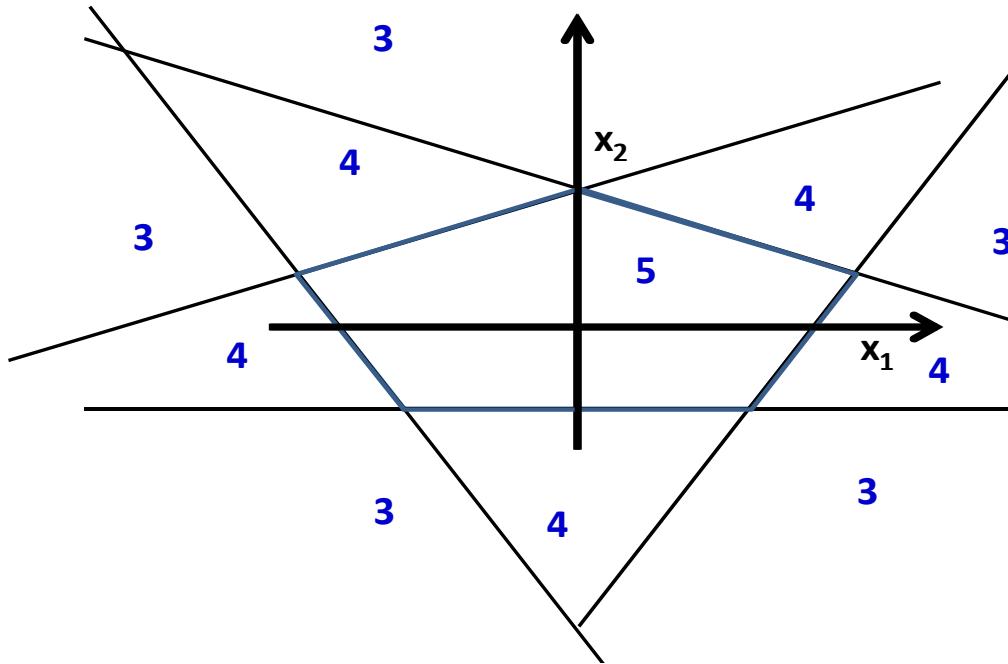
Combining Gates to build MLP by hand



19

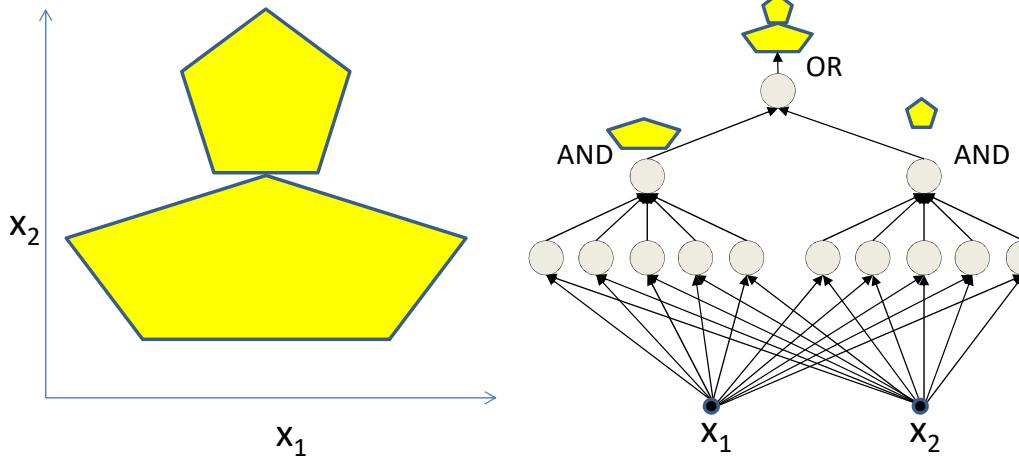
Source: 11785 lecture notes

More complex Decision Boundaries



Source: 11785 lecture notes

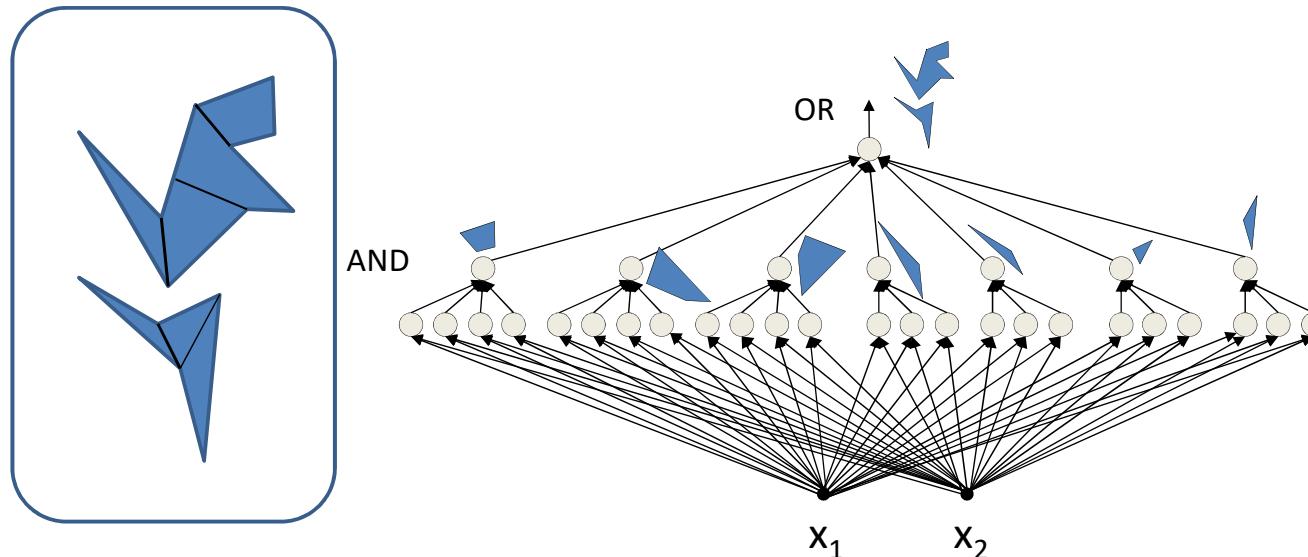
More complex Decision Boundaries



- Network to fire if the input is in the yellow area
 - “OR” two polygons
 - A third layer is required

Source: 11785 lecture notes

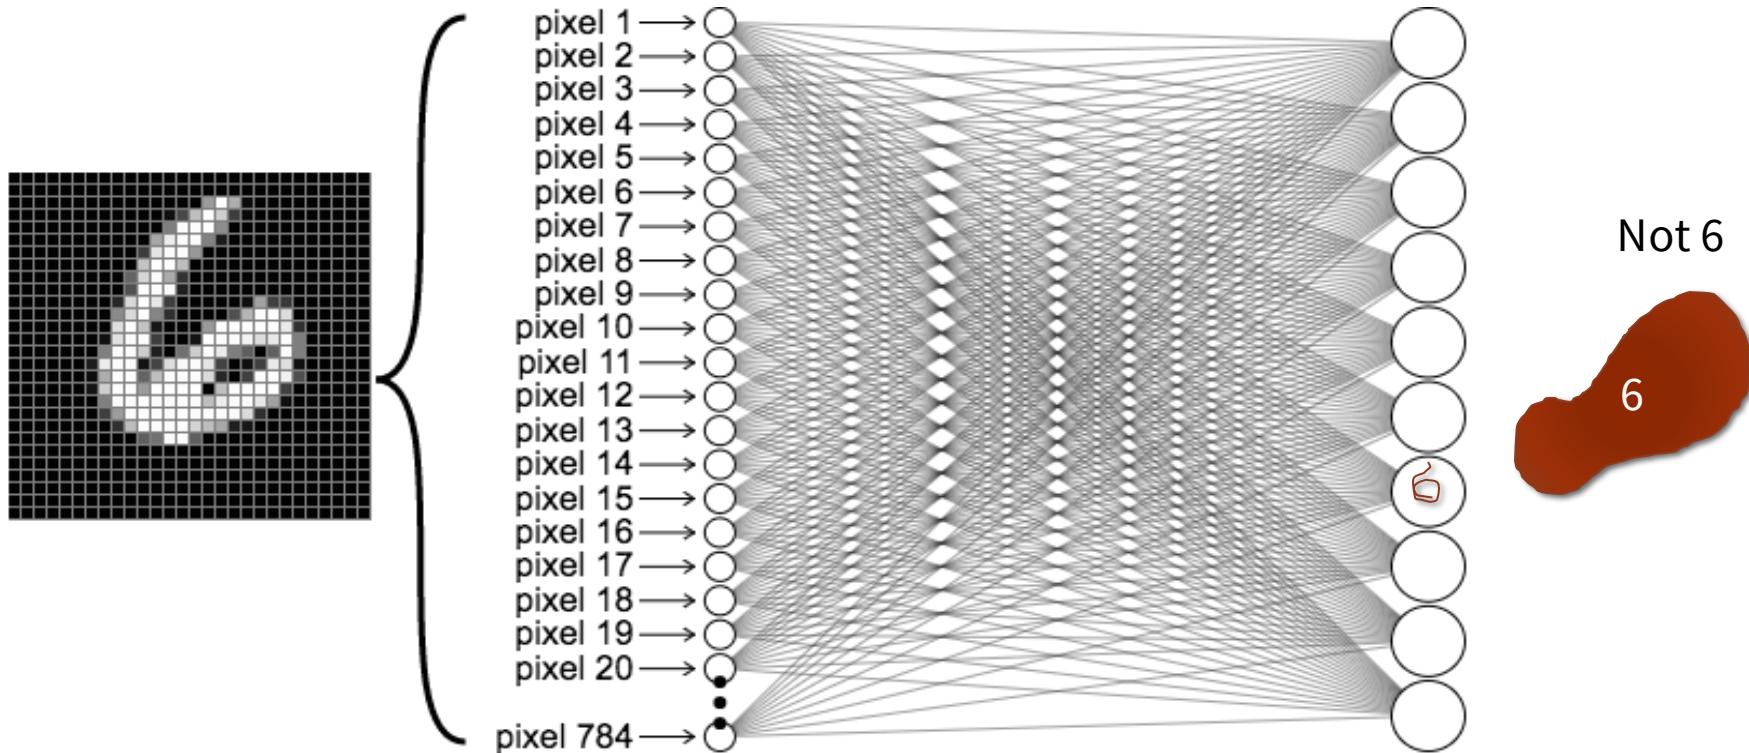
More complex Decision Boundaries



- Can compose *arbitrarily* complex decision boundaries

Source: 11785 lecture notes

Digit Recognition



Course Topics

- Introduction to Deep Learning and its application
- Neural Networks
- Convolutional Neural Networks (CNN)
- Training and Testing CNN
- Interpretability of Deep Learning
- Solving Engineering problems using Deep Learning



Pytorch

```
# loading PyTorch
import torch

# cuda
import torch.cuda as tCuda # various functions and settings
torch.backends.cudnn.deterministic = True # deterministic ML?
torch.backends.cudnn.benchmark = False # deterministic ML?
torch.cuda.is_available # check if cuda is is_available
tensor.cuda() # moving tensor to gpu
tensor.cpu() # moving tensor to cpu
tensor.to(device) # copy tensor to device xyz
torch.device('cuda') # or 'cuda0', 'cuda1' if multiple devices
torch.device('cpu') # default

# static computation graph/C++ export preparation
torch.jit.trace()
from torch.jit import script, trace
@script

# load and save a model
torch.save(model, 'model_file')
model = torch.load('model_file')
model.eval() # set to inference
torch.save(model.state_dict(), 'model_file') # only state dict
model = ModelClass()
model.load_state_dict(torch.load('model_file'))
```